

Types of models to train

Your final submission should include single model. The model set you should try to come up with best model per type of model:

1. Identify best model from: H2O-3 GLM - try different combinations of regularization

Evaluation metric: AUCPR

Feature engineering

You should train/fit categorical features scalers and encoders on Train only. Use `transform` or equivalent function on Validation/Test datasets.

It is important to understand all the steps before model training, so that you can reliably replicate and test them to produce scoring function.

You should generate various new features. Examples of such features can be seen in the Module-3 lecture on GLMs.

Your final model should have at least **10** new engineered features.

On-hot-encoding, label encoding, and target encoding **is not included in the 10** features to create.

You can attempt target encoding, however the technique is not expected to produce improvement for Linear models.

Ideas for Feature engineering for various types of variables:

1. <https://docs.h2o.ai/driverless-ai/1-10-lts/docs/userguide/transformations.html>
2. GLM lecture and hands-on (Module-3)

Note:

- You don't have to perform feature engineering using H2O-3 even if you decided to use H2O-3 GLM for model training.
- It is OK to perform feature engineering using any technique, as long as you can replicate it correctly in the Scoring function.

Threshold calculation

You will need to calculate optimal threshold for class assignment using F1 metric:

- Using H2O-3, use F1

You will need to find optimal probability threshold for class assignment, the threshold that maximizes above F1.

H2O Model

```
import h2o
try:
    h2o.cluster().shutdown()
```

```

except:
    pass

H2O session _sid_a34c closed.

from h2o.frame import H2OFrame
h2o.init(max_mem_size = "4G", nthreads=16)

Checking whether there is an H2O instance running at
http://localhost:54321..... not found.
Attempting to start a local H2O server...
; Java HotSpot(TM) 64-Bit Server VM (build 25.361-b09, mixed mode)
Starting server from D:\Work\Gre\UTD\Courses\Fall\MIS6341\Softwares\
Python\ml-fall-2023\Lib\site-packages\h2o\backend\bin\h2o.jar
Ice root: C:\Users\Asus\AppData\Local\Temp\tmpb9pjpbvq
JVM stdout: C:\Users\Asus\AppData\Local\Temp\tmpb9pjpbvq\
h2o_Asus_started_from_python.out
JVM stderr: C:\Users\Asus\AppData\Local\Temp\tmpb9pjpbvq\
h2o_Asus_started_from_python.err
Server is running at http://127.0.0.1:54321
Connecting to H2O server at http://127.0.0.1:54321 ... successful.

-----
H2O_cluster_uptime:      08 secs
H2O_cluster_timezone:   America/Chicago
H2O_data_parsing_timezone: UTC
H2O_cluster_version:    3.42.0.3
H2O_cluster_version_age: 2 months and 14 days
H2O_cluster_name:       H2O_from_python_Asus_zbkukp
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 3.556 Gb
H2O_cluster_total_cores: 16
H2O_cluster_allowed_cores: 16
H2O_cluster_status:     locked, healthy
H2O_connection_url:     http://127.0.0.1:54321
H2O_connection_proxy:   {"http": null, "https": null}
H2O_internal_security:  False
Python_version:         3.10.11 final
-----

```

We import an H2O DataFrame (df_h) from a CSV file located at the specified path and then display its first few rows using df_h.head(). It is imported like a pandas data frame, only difference being that it resides inside a cluster

```

from h2o.estimators.glm import H2OGeneralizedLinearEstimator

# Specify the path to your CSV file
csv_file_path =
'D:/Work/Gre/UTD/Courses/Fall/MIS6341/Softwares/Python/ml-fall-2023/
Project1/h2o_df.csv'

```

```
df_h = h2o.import_file(csv_file_path)
```

```
df_h.head()
```

```
Parse progress: |
```



```
(done) 100%
```

City	State	Zip	Bank	BankState	NAICS
NoEmp	NewExist	CreateJob	RetainedJob	FranchiseCode	
UrbanRural	RevLineCr	LowDoc	DisbursementGross		
BalanceGross	GrAppv	SBA_Appv	MIS_Status		
Log_DisbursementGross	Log_GrAppv	Log_SBA_Appv			
Log_BalanceGross	TotalJobs	IncomeToLoanRatio			
EmployeesToLoanRatio	JobPerLoan	Gauren_SBA_Appv	DefaultRate		
0.11465	0.184773	93001	0.0314465	0.218517	235910
0.600407	0.17044	-0.0353733	-0.0454543		1
0.0716743	0.15307	0.187063		0.358949	-0.00229552
0.394801	0.410973		0	0.306712	
0.332752	0.344279		-0.00229816	-0.0808276	
0.873414		1.46094	-0.196674		0.960651
17.5096					
0.137597	0.165992	44039	0.128698	0.159167	484121 -0.1536
0.17044	-0.0353733		-0.0454543		0 0.243491
0.15307	0.187063		-0.614207	-0.00229552	-0.594552
-0.600302		0		-0.952455	-0.902762 -
0.917047		-0.00229816	-0.0808276		1.02316
0.255872	0.134645		0.990421		17.5096
0.139151	0.116799	68122	0.175694	0.159167	451120
0.115688	0.17044		-0.0139777	0.0185835	38510
0.243491	0.15307	0.187063		2.31704	-0.00229552
2.752	2.53323		0		1.19907
1.32229	1.26221		-0.00229816	0.00460579	
0.914656		0.045668	0.00181815		1.08636
17.5096					
0.140704	0.197662	14208	0.118949	0.167297	321114 0.51962
0.17044	-0.0353733		-0.0454543		1 0.0716743
0.15307	0.187063		-0.0046858	-0.00229552	0.0251141
0.00123195		0		-0.00469681	0.0248039
0.00123119		-0.00229816	-0.0808276		-3.80356
421.787	-65.6095		20.3856		17.5096
0.140354	0.144273	16335	0.193277	0.0783071	0

```

0.102223    0.17044    -0.0353733    -0.0454543    1
0.0716743    0.15307    0.187063    0.584056    -0.00229552
0.623655    0.637638    0    0.459989
0.48468    0.493255    -0.00229816    -0.0808276
0.915969    0.160316    -0.126761    0.978071
17.5096
0.164738    0.188249    77381    0.141056    0.180023    532490    -
0.140136    0.17044    -0.0268151    -0.0411851    1
0.243491    0.251569    0.187063    -0.607811    -0.00229552
-0.608635    -0.60902    0    -0.93601    -
0.938115    -0.939099    -0.00229816    -0.0680002
0.998014    0.230101    0.111655    0.999368
17.5096
0.172811    0.223814    48334    0.214858    0.198182    811111    -
0.0458851    0.17044    -0.0353733    -0.0113008    0
0.243491    0.251569    0.187063    -0.584843    -0.00229552
-0.65335    -0.636699    0    -0.879098    -
1.05944    -1.01252    -0.00229816    -0.0466741
0.918554    0.0720672    0.0733063    1.02615
17.5096
0.275041    0.184773    90016    0.182257    0.180023    453210    -
0.0997428    0.17044    -0.0353733    -0.0454543    50564
0.0716743    0.15307    0.0897581    -0.437584    -0.00229552
-0.414989    -0.391073    1    -0.575514    -
0.536125    -0.496057    -0.00229816    -0.0808276
1.11893    0.255049    0.206682    1.06116
17.5096
0.170819    0.179137    19805    0.214858    0.198182    722110    -
0.113207    0.17044    -0.0353733    -0.0326467    1
0.243491    0.251569    0.187063    -0.680008    -0.00229552
-0.661447    -0.641712    0    -1.13946    -
1.08308    -1.02642    -0.00229816    -0.06802
1.05968    0.176414    0.105998    1.03075
17.5096
0.0820406    0.0793893    85267    0.329864    0.0762233    484220    -
0.140136    0.186978    -0.0310942    -0.0369159    1
0.187265    0.15307    0.187063    -0.664423    -0.00229552
-0.645604    -0.631905    0    -1.0919    -
1.03734    -0.999413    -0.00229816    -0.0680101
1.05146    0.221768    0.107627    1.02168
17.5096
[10 rows x 29 columns]

```

```
df_h.describe()
```

```
Rows:566472
```

```
Cols:29
```

Bank	City	BankState	State	NAICS	Zip	NoEmp
------	------	-----------	-------	-------	-----	-------

[illegible]

40.51182843681707	92006.0	0.2434914068009948
0.25156900049734354	0.18706258164639383	37.39775945774547
552.889522851626	18.58692184996103	23.199528251630852
1.0	3.6479991104368574	2.9748620909714023
3.186333139238373	6.316965249625691	75.14418220289946
6396.558729319238	37035.12034202456	11627.399763221902
1043.390760820282	17.509603299015662	
sigma	0.06639769763247672	0.04199181924329759
0.12227701398294173	0.07450052965800262	31185.22931842036
1.000000882657328	0.007437873600660589	263321.75739069114
1.0000008826573283	12747.306218543337	1.0000008826573281
0.041041501003144853	0.031956251232178115	0.07945017170394587
1.00000088265733	1.0000008826573281	1.0000008826573277
0.38004955887207637	0.6751342455877893	1.0000008826573283
0.6724602419470982	0.015557287489121684	0.673495391264487
48.40854150735596	57.47035513066714	1.9967248008580234
11.623336619659533	4.396234168225107e-15	29.416621353104784
zeros	0	0
20864	0	176
0	0	127266
131369	0	0
0	0	0
0	467285	0
0	0	0
0	0	0
0	0	0
missing	0	0
0	0	0
0	0	0
0	0	0
0	0	0
0	0	0
0	0	0
0	0	0
0	0	0
0	0	0
0	0.11464974937637376	0.18477258385515263
0.031446540880519576	0.21851689918783204	93001.0
0.6004066600002538	0.17044022221184818	235910.0
0.04545428677284537	1.0	-0.03537327806289165
0.1530699251248866	0.18706258164639383	0.07167428097535124
0.002295517679303139	0.394801069280922	0.35894908638142187
0.0	0.3067116704358818	-
0.3442792686074715	-0.002298156418966926	0.4109726208502541
0.873413624583551	1.4609407769259248	0.3327518024385963
0.9606505378974516	17.509603299015662	-0.08082756483573703
1	0.137596588115639	-0.19667384330497317
0.12869797760320908	0.15916691590309687	44039.0
0.153600434175742	0.17044022221184818	484121.0
0.04545428677284537	0.0	-0.03537327806289165
		-
		0.2434914068009948

0.1530699251248866	0.18706258164639383	-0.6142071281443656	-
0.002295517679303139	-0.5945517349325695	-0.6003022273404686	
0.0	-0.9524546549496404	-0.9027619966303708	-
0.9170465858109552	-0.002298156418966926	-0.08082756483573703	
1.0231631671025114	0.25587183785114576	0.1346447858336776	
0.99042067121194	17.509603299015662		
2	0.1391509433962264	0.11679920477137178	68122.0
0.1756944756092088	0.15916691590309687	451120.0	
0.11568781374425648	0.17044022221184818	-0.013977733515888482	
0.01858352805387556	38510.0	0.2434914068009948	
0.1530699251248866	0.18706258164639383	2.3170363593098213	-
0.002295517679303139	2.751995739818049	2.533233145474642	
0.0	1.199071721409156	1.3222878957003827	
1.2622133569935532	-0.002298156418966926	0.00460579453798708	
0.9146557881768466	0.045668048340090896	0.0018181486951624856	
1.0863570708974033	17.509603299015662		
3	0.1407035175879397	0.19766159276417383	14208.0
0.1189488243430152	0.16729655837889998	321114.0	
0.5196201856242543	0.17044022221184818	-0.03537327806289165	-
0.04545428677284537	1.0	0.07167428097535124	
0.1530699251248866	0.18706258164639383	-0.004685798049922222	-
0.002295517679303139	0.02511407838264226	0.0012319496005647338	
0.0	-0.0046968108174749315	0.024803902380498698	
0.00123119137332351	-0.002298156418966926	-0.08082756483573703	
-3.803563106619152	421.7868859132362	-65.60947363324371	
20.385637830581544	17.509603299015662		
4	0.14035375121577756	0.14427252985884909	16335.0
0.19327731092436976	0.07830707560361704	0.0	
0.10222340134825655	0.17044022221184818	-0.03537327806289165	-
0.04545428677284537	1.0	0.07167428097535124	
0.1530699251248866	0.18706258164639383	0.5840563957913015	-
0.002295517679303139	0.6236549207893809	0.6376376730309333	
0.0	0.4599888961620474	0.4846797319622258	
0.4932547601265484	-0.002298156418966926	-0.08082756483573703	
0.91596908478613	0.16031581205412476	-0.12676096199199305	
0.9780710067284966	17.509603299015662		
5	0.16473815924050095	0.1882487805974511	77381.0
0.1410558507971523	0.18002315611603092	532490.0	-
0.14013602177974208	0.17044022221184818	-0.02681506024409038	-
0.041185099117730634	1.0	0.2434914068009948	
0.25156900049734354	0.18706258164639383	-0.6078106173677496	-
0.002295517679303139	-0.6086350488715516	-0.6090201139628024	
0.0	-0.936010436896243	-0.9381147754788066	-
0.9390991626800832	-0.002298156418966926	-0.06800015936182102	
0.998014028490483	0.23010081041149533	0.11165503043791804	
0.9993677300922869	17.509603299015662		
6	0.1728110599078341	0.22381434467720585	48334.0
0.21485771162285663	0.1981815002622836	811111.0	-
0.045885135007742606	0.17044022221184818	-0.03537327806289165	-

```

0.011300785531927536      0.0      0.2434914068009948
0.25156900049734354      0.18706258164639383      -0.5848427454291907      -
0.002295517679303139      -0.6533495706278197      -0.6366994039887123
0.0      -0.8790979038493719      -1.059438414961258      -
1.012524699207778      -0.002298156418966926      -0.04667406359481919
0.9185539389001204      0.07206718699638688      0.07330627813128383
1.026150749529212      17.509603299015662
7      0.27504127682993945      0.18477258385515263      90016.0
0.182256711409396      0.18002315611603092      453210.0      -
0.0997427845917423      0.17044022221184818      -0.03537327806289165      -
0.04545428677284537      50564.0      0.07167428097535124
0.1530699251248866      0.08975812818488481      -0.4375844699919985      -
0.002295517679303139      -0.4149894822105479      -0.391072948404457
1.0      -0.5755143250547002      -0.5361254527846345      -
0.4960568023622707      -0.002298156418966926      -0.08082756483573703
1.1189331089693224      0.25504905158662605      0.20668155433789612
1.0611561958035405      17.509603299015662
8      0.1708185053380783      0.17913669064748203      19805.0
0.21485771162285663      0.1981815002622836      722110.0      -
0.11320719698774225      0.17044022221184818      -0.03537327806289165      -
0.03264672380750118      1.0      0.2434914068009948
0.25156900049734354      0.18706258164639383      -0.6800077262795612      -
0.002295517679303139      -0.6614474761427344      -0.6417121887965542
0.0      -1.1394584281034796      -1.0830760321756987      -
1.0264186736605694      -0.002298156418966926      -0.06802000187039284
1.0596771234076527      0.17641428503960202      0.10599767786545444
1.0307541101614277      17.509603299015662
9      0.0820405814622732      0.07938931297709924      85267.0
0.32986399789891396      0.07622333751568382      484220.0      -
0.14013602177974208      0.1869783599731616      -0.031094169153491016      -
0.03691591146261591      1.0      0.18726450640542577
0.1530699251248866      0.18706258164639383      -0.6644233740896465      -
0.002295517679303139      -0.6456037479613795      -0.6319045663464287
0.0      -1.091904955472731      -1.0373396358189255      -
0.9994130438532118      -0.002298156418966926      -0.06801008061610692
1.0514615805535896      0.2217676991796185      0.10762713903039244
1.0216791938918204      17.509603299015662
[566472 rows x 29 columns]

```

To check if the data frame is indeed H2OFrame

```

if isinstance(df_h, h2o.H2OFrame):
    print('It is H2O data frame')
else:
    print('It is not H2O data frame')

It is H2O data frame

```


The provided code splits an H2O DataFrame into training, validation, and test sets and separates predictor columns from the response column for machine learning tasks.

The first number, 0.7, specifies that 70% of the data will be used for training (the "train" subset).

The second number, 0.15, specifies that 15% of the data will be used for validation (the "valid" subset).

The remaining 15% not specified is implicitly used for testing (the "test" subset).

The `seed=1234` parameter is used to set the random seed for reproducibility, ensuring that the same split is obtained when the code is executed multiple times.

```
# Split the data as described above
train, valid, test = df_h.split_frame([0.7, 0.15], seed=1234)

# Prepare predictors and response columns
train_X = df_h.columns
train_y = "MIS_Status"
train_X.remove(train_y)
```

Since we already imported the H2O GLM estimator, we will just instantiate our model. For simplicity, the name of our model will be `glm`. To build a GLM, you just need to define the family, and you are ready to go. However, we will set a random seed for reproducibility purposes, and also a model id to be able to retrieve the model later on if we need to access it. You can instantiate your GLM, as shown below.

```
glm = H2OGeneralizedLinearEstimator(family = "binomial", seed = 42,  
model_id = 'default_glm')
```

```
%time glm.train(x = train_X, y = train_y, training_frame = train,
validation_frame = valid)
```

```
glm Model Build progress: |
```

```
| (done) 100%
```

```
CPU times: total: 15.6 ms
```

Wall time: 1.73 s

Model Details

=====

H20GeneralizedLinearEstimator : Generalized Linear Modeling

Model Key: default glm

GLM Model: summary

family	link	regularization
--------	------	----------------

number of predictors total number of active predictors

number of iterations training frame

```

-----
      binomial  logit  Elastic Net (alpha = 0.5, lambda = 2.505E-4 )
27                                     22                                     4
py_17_sid_b6e5

```

ModelMetricsBinomialGLM: glm
 ** Reported on train data. **

MSE: 0.12059947193887977
 RMSE: 0.34727434679066027
 LogLoss: 0.38760076479874667
 AUC: 0.7768657635596997
 AUCPR: 0.4455013625622545
 Gini: 0.5537315271193994
 Null degrees of freedom: 396643
 Residual degrees of freedom: 396621
 Null deviance: 367114.9304608208
 Residual deviance: 307479.03550566814
 AIC: 307525.03550566814

Confusion Matrix (Act/Pred) for max f1 @ threshold = 0.22939337808642146

	0	1	Error	Rate
0	267868	59606	0.182	(59606.0/327474.0)
1	29773	39397	0.4304	(29773.0/69170.0)
Total	297641	99003	0.2253	(89379.0/396644.0)

Maximum Metrics: Maximum metrics at their respective thresholds

metric	threshold	value	idx
--------	-----------	-------	-----

max f1	0.229393	0.468529	225
max f2	0.121245	0.596597	301
max f0point5	0.345963	0.462715	165
max accuracy	0.509815	0.83613	100
max precision	0.940776	0.943396	2
max recall	0.00176269	1	399
max specificity	0.972444	0.999997	0
max absolute_mcc	0.241421	0.341047	218
max min_per_class_accuracy	0.171115	0.703744	264
max mean_per_class_accuracy	0.168308	0.706072	266
max tns	0.972444	327473	0
max fns	0.972444	69159	0
max fps	0.00176269	327474	399
max tps	0.00176269	69170	399
max tnr	0.972444	0.999997	0
max fnr	0.972444	0.999841	0
max fpr	0.00176269	1	399
max tpr	0.00176269	1	399

Gains/Lift Table: Avg response rate: 17.44 %, avg score: 17.44 %

group	cumulative_data_fraction	lower_threshold	lift
cumulative_lift	response_rate	score	
cumulative_response_rate	cumulative_score	capture_rate	
cumulative_capture_rate	gain	cumulative_gain	
kolmogorov_smirnov			
-----	-----	-----	-----
-----	-----	-----	-----
-----	-----	-----	-----
-----	-----	-----	-----
1	0.0100014	0.734962	4.23823
4.23823	0.739098	0.801263	0.739098
0.801263	0.0423883	0.0423883	323.823
323.823	0.0392278		
2	0.0200003	0.654405	3.70722
3.97276	0.646495	0.693539	0.692802
0.747407	0.0370681	0.0794564	270.722
297.276	0.0720146		
3	0.0300017	0.596196	3.30877
3.75141	0.57701	0.62454	0.654202
0.706448	0.0330924	0.112549	230.877
275.141	0.0999829		
4	0.0400006	0.547301	3.15779
3.60302	0.550681	0.570568	0.628325
0.672482	0.0315744	0.144123	215.779
260.302	0.126116		
5	0.0500002	0.506656	2.94017
3.47044	0.51273	0.52673	0.605203
0.643329	0.0294058	0.173529	194.017
247.044	0.149619		
6	0.1000002	0.387661	2.54795
3.00921	0.444332	0.440826	0.52477
0.54208	0.127396	0.300925	154.795
200.921	0.243363		
7	0.1500001	0.314515	2.14662
2.72168	0.374344	0.348613	0.474629
0.477592	0.10733	0.408255	114.662
172.168	0.312803		
8	0.2000001	0.264456	1.76581
2.48272	0.307937	0.28817	0.432956
0.430237	0.0882897	0.496545	76.5812
148.272	0.359181		
9	0.2999999	0.199768	1.3652
2.11022	0.238075	0.229392	0.367996
0.363289	0.136519	0.633063	36.5201
111.022	0.403415		
10	0.4000001	0.156437	1.06143
1.84801	0.1851	0.177086	0.322272

0.316738		0.106144	0.739208	6.14268
84.8015		0.410855		
11	0.5		0.122838	0.79862
1.63814		0.13927	0.13907	0.285672
0.281205		0.0798612	0.819069	-20.138
63.8138		0.386464		
12	0.599999		0.0957864	0.610386
1.46685		0.106444	0.108834	0.255801
0.252476		0.061038	0.880107	-
38.9614	46.6847		0.339273	
13	0.700001		0.0740264	0.484596
1.32652		0.0845078	0.0845272	0.23133
0.228483		0.0484603	0.928567	-
51.5404	32.6524		0.276845	
14	0.799999		0.0569197	0.357239
1.20536		0.0622983	0.0652598	0.210201
0.208081		0.0357236	0.964291	-
64.2761	20.5364		0.198993	
15	0.899998		0.0410503	0.249099
1.09911		0.0434399	0.0490687	0.191672
0.190413		0.0249096	0.989201	-
75.0901	9.91135		0.108044	
16	1		4.27023e-11	0.107993 1
0.0188327		0.0301829	0.174388	0.17439
0.0107995	1		-89.2007 0	
0				

ModelMetricsBinomialGLM: glm
 ** Reported on validation data. **

MSE: 0.12231119069488473
 RMSE: 0.3497301684082812
 LogLoss: 0.3920168718665924
 AUC: 0.7738999913401364
 AUCPR: 0.43815403112953155
 Gini: 0.5477999826802729
 Null degrees of freedom: 84950
 Residual degrees of freedom: 84928
 Null deviance: 79045.0523443093
 Residual deviance: 66604.45056387779
 AIC: 66650.45056387779

Confusion Matrix (Act/Pred) for max f1 @ threshold =
 0.2239578972008154

	0	1	Error	Rate
0	56619	13383	0.1912	(13383.0/70002.0)
1	6390	8559	0.4275	(6390.0/14949.0)
Total	63009	21942	0.2328	(19773.0/84951.0)

5	0.0500053		0.508851	2.86479	
3.3885		0.504122	0.527854	0.596281	
0.643937		0.0286307	0.169443		186.479
238.85		0.144943			
6	0.100011		0.388025	2.53502	
2.96176		0.446092	0.44178	0.521186	
0.542859		0.126764	0.296207		153.502
196.176		0.238094			
7	0.150004		0.314092	2.0673	
2.66365		0.363786	0.348682	0.468728	
0.478143		0.103351	0.399558		106.73
166.365		0.302847			
8	0.200009		0.265068	1.69893	
2.42246		0.298964	0.288529	0.426285	
0.430737		0.0849555	0.484514		69.893
142.246		0.345261			
9	0.300008		0.199691	1.41616	
2.08704		0.249205	0.229498	0.36726	
0.36366		0.141615	0.626129		41.6165
108.704		0.395764			
10	0.400007		0.156012	1.07968	
1.83521		0.189994	0.176796	0.322945	
0.316945		0.107967	0.734096		7.96836
83.5207		0.405434			
11	0.500006		0.122342	0.814111	
1.63099		0.143261	0.13849	0.287009	
0.281255		0.0814101	0.815506		-
18.5889	63.0993	0.382876			
12	0.600005		0.0954449	0.661591	
1.46943		0.116421	0.108555	0.258578	
0.252472		0.0661583	0.881664		-
33.8409	46.9429	0.341808			
13	0.700004		0.0737534	0.482981	
1.32851		0.0849912	0.0840837	0.233781	
0.228417		0.0482975	0.929962		-
51.7019	32.851	0.279066			
14	0.800002		0.0567782	0.346516	
1.20576		0.060977	0.0650276	0.212181	
0.207994		0.0346511	0.964613		-
65.3484	20.5763	0.199763			
15	0.900001		0.0411062	0.253532	
1.09996		0.0446145	0.0491226	0.193562	
0.190342		0.0253529	0.989966		-
74.6468	9.99607	0.109177			
16	1		6.97019e-06	0.100342	1
0.0176574		0.0301928	0.175972	0.174327	
0.0100341	1		-89.9658	0	
0					

Scoring History:

timestamp	duration	iterations	negative_log_likelihood	objective	training_rmse
training_logloss	training_r2	training_auc	training_pr_auc	training_lift	training_classification_error
validation_rmse	validation_logloss	validation_r2	validation_auc	validation_pr_auc	validation_lift
validation_classification_error					
2023-11-05 20:51:48	0.000 sec	0	183557		
0.462776					
2023-11-05 20:51:48	0.214 sec	1	156657		
0.395782					
2023-11-05 20:51:48	0.285 sec	2	153810		
0.388709					
2023-11-05 20:51:48	0.340 sec	3	153734		
0.38859					
2023-11-05 20:51:48	0.403 sec	4	153740		
0.388591					
0.34727434679066027			0.38760076479874667		
0.16236930409227046			0.7768657635596997		0.4455013625622545
4.238233490486189			0.22533808654612197		0.3497301684082812
0.3920168718665924			0.15650872440673047		0.7738999913401364
0.43815403112953155			4.098244856039948		0.2327577073842568

Variable Importances:

variable	relative_importance	scaled_importance
percentage		
Log_DisbursementGross	1.0742710828781128	1.0
0.17945711926553892		
GrAppv	0.7878198623657227	0.7333529450081168
0.13160540691605582		
Log_GrAppv	0.7507895827293396	0.6988828003429787
0.12541949405378375		
DisbursementGross	0.7272646427154541	0.6769842866541813
0.12148964987099521		
Log_SBA_Appv	0.5599473714828491	0.5212347054736658
0.0935392787055257		
Bank	0.5510706901550293	0.5129717246773866
0.09205642797527898		
City	0.3853748142719269	0.3587314416389738
0.06437691110650404		

```

UrbanRural          0.3724907636642456      0.34673814607975306
0.06222462883494611
SBA_Appv            0.1293647140264511      0.12042092176572985
0.021610391719378712
RevLineCr           0.11934459209442139      0.11109355356999986
0.019936529092644
---
---
BankState           0.03160914033651352      0.029423802651215834
0.005280310861624921
FranchiseCode       0.025763940066099167      0.023982717655467556
0.004303869422608988
Gauren_SBA_Appv     0.008871283382177353      0.008257956044399915
0.0014819490027494534
Log_BalanceGross    0.004079081583768129      0.0037970691464948828
0.0006814110906820303
JobPerLoan          5.3359181038104e-05      4.967012692471231e-05
8.913657891462553e-06
RetainedJob         0.0                          0.0
0.0
BalanceGross        0.0                          0.0
0.0
TotalJobs           0.0                          0.0
0.0
IncomeToLoanRatio   0.0                          0.0
0.0
EmployeesToLoanRatio 0.0                          0.0
0.0
[27 rows x 4 columns]

[tips]
Use `model.explain()` to inspect the model.
--
Use `h2o.display.toggle_user_tips()` to switch on/off this section.

```

From the summary results, we can see the GLM performance. We will focus on the Area Under the Curve (AUC), and since we have a very imbalanced dataset, we will be looking at the F1 score. Additionally, we will also take a quick look at the misclassification error and logloss.

From the report, we can look at the metrics on the training and validation data, and we see that the training AUCPR was AUCPR: 0.4455013625622545 while the validation AUCPR: 0.43815403112953155

``AUCPR (Area Under the Precision-Recall Curve) is a metric that measures the trade-off between precision and recall for a binary classification model.

AUCPR summarizes how well a model is at ranking the positive class instances higher than the negative class instances.

Precision is a measure of the model's ability to correctly identify positive instances among the instances it predicts as positive. It's the ratio of true positive predictions to the total positive predictions

Recall (or Sensitivity) is a measure of the model's ability to identify all the positive instances correctly. It's the ratio of true positive predictions to the total actual positive instances.

The reported values represent the AUCPR score for your model on both the training and validation datasets:

...

Training AUCPR: 0.4455013625622545> This indicates that on the training dataset, the model achieved an AUCPR of approximately 0.4455. This score reflects how well the model performs on the data it was trained on. It suggests that the model is reasonably effective at ranking and identifying positive instances relative to negative instances within the training dataset.

<Validation AUCPR: 0.43815403112953155> This indicates that on the validation dataset, the model achieved an AUCPR of approximately 0.4382. The validation dataset is a separate dataset that the model did not see during training. The AUCPR score on the validation dataset measures the model's generalization performance. While slightly lower than the training AUCPR, this score still suggests that the model maintains its ability to rank positive instances effectively when applied to new, unseen data.

The goal is to have a model that generalizes well, which means it performs consistently on both the training and validation datasets. In this case, the AUCPR scores are reasonably close, indicating that the model is performing consistently, and it is not exhibiting significant overfitting (a situation where the model fits the training data too closely but performs poorly on new data).

From the report, we can also see the max F1 score as well as all the metrics for our model with their respective thresholds. For the default GLM, we obtained a training F1 score of 0.2293934 and a validation F1 score of 0.2239579.

F1 Score: The F1 score is a metric that combines both precision and recall into a single value, providing a balanced measure of a model's performance. It is especially useful when dealing with imbalanced datasets

Training F1 Score: This is the F1 score calculated on the training dataset. It measures how well the model performs on the data it was trained on. A training F1 score of 0.2293934 means that the model achieved an F1 score of approximately 0.2294 when tested on the training data.


Validation F1 Score: This is the F1 score calculated on the validation dataset. The validation dataset is a separate set of data that the model did not see during training. The validation F1 score of 0.2239579 indicates that when the model is applied to new, unseen data (the validation dataset), it achieved an F1 score of approximately 0.2240. This demonstrates how well the model generalizes to data it has not been trained on.

Thresholds: In binary classification, different threshold values can be used to determine whether a prediction is classified as the positive or negative class. The explanation suggests that the F1 scores reported are associated with specific threshold values. Different thresholds can affect the trade-off between precision and recall, and they are used to fine-tune the model's performance.

In summary, the explanation indicates that a machine learning model was evaluated using F1 scores on both the training and validation datasets. The training F1 score represents its performance on the training data, while the validation F1 score reflects how well the model generalizes to new, unseen data. The thresholds mentioned suggest that different threshold values were applied to calculate these scores, influencing the precision and recall trade-off.

Plot the Scoring history for any of our models, as shown below:

```
glm.plot(metric='negative_log_likelihood')
```




```
<h2o.plot._plot_result._MObject at 0x24155b7a260>
```

We can also generate a variable importance plot to see how each of our features contribute to the linear model.

As we can see after 2 iterations, the scores don't really improve after this time. We can also use the default number of iterations and use early stopping; that way, the model will stop training when it is no

longer improving. We will use early stopping when we start tuning our models.

```
glm.varimp_plot()
```



```
<h2o.plot._plot_result._MObject at 0x241a2eeb430>
```

```
<Figure size 640x480 with 0 Axes>
```

From the variable importance plot, we can see that the most significant feature is Log_DisbursementGross. We can also see Gr_Appv, (Log_Gr_Appv) DisbursementGross, and Log_SBA_Appv are the next most important variables. As this is understood by the fact, that if your

DisbursementGross provides information about the size of the loans granted to small businesses. This information is crucial for understanding the financial impact of SBA loans on the businesses they support. The loan amount disbursed is often indicative of the level of risk associated with a borrower. Larger loan amounts may indicate higher financial stability, while smaller loans may be associated with smaller or riskier businesses.

SBA_Appv (Log_SBA_Appv) The variable confirms that the SBA has approved a loan for a specific borrower. It signifies that the borrower has successfully gone through the SBA's application and

approval process Lenders and borrowers use this variable to determine if they are eligible for SBA loans and to understand the maximum loan amount that can be approved for their business. The variable is important for assessing the economic impact of SBA loans. By analyzing the approved loan amounts, one can estimate the potential economic impact in terms of job creation, business expansion, and overall economic growth.

he "Bank" variable helps identify the specific financial institutions that are participating in the SBA loan program. This information is crucial for understanding which banks are actively providing SBA loans to small businesses.

The Top 4 Variables with their relative importance is as follows:

Log_DisbursementGross: 1.0742711

Log_Gr_Appv: 0.7507896

Log_SBA_Appv: 0.5599474

Bank: 0.5510707

```
glm.predict(valid).head(10)
```

```
glm prediction progress: |
```

```
| (done) 100%
```

```

predict      p0      p1
-----
0  0.970777  0.029223
1  0.471275  0.528725
0  0.951736  0.0482639
0  0.951392  0.0486081
0  0.945459  0.0545409
0  0.92252   0.0774796
0  0.855034  0.144966
0  0.865822  0.134178
0  0.919617  0.0803829
1  0.759493  0.240507
[10 rows x 3 columns]

```

These columns contain the predicted probabilities for each class. "p0" represents the probability of an observation belonging to class 0 (Not Defaulted), and "p1" represents the probability of it belonging to class 1 (Defaulted on Loan). These probabilities can be used to assess the model's confidence in its predictions.

H2O Model Tuning

We Tune our model with `lambda_search = True`, as this will automatically tune the model. Other parameters that we can alter are `max_active_predictors` (feature selection parameter), `n_lambdas`, which allows you to specify the number of lambda values, or the regularization strengths, to be used in the elastic net regularization path, and `solver`, which specify the algorithm or optimization method that the GLM model should use to find the solution

A value of $\alpha = 1$ represents Lasso Regularization and a value of $\alpha = 0$ produces Ridge regression

lambda is employed for regularization strength
missing_value_handling
parameter allows to specify how we want to handle any missing data
(options are skip and MeanImputation)

```
glm_grid = h2o.grid.H2OGridSearch (
    H2OGeneralizedLinearEstimator(family = "binomial",
                                   lambda_search = True),
    hyper_params = {"alpha": [x*0.01 for x in range(0, 50)],
                    "missing_values_handling" : ["Skip",
                                                  "MeanImputation"]},
    grid_id = "glm_random_grid",
    search_criteria = {
        "strategy":"RandomDiscrete",
        "max_models":300,
        "max_runtime_secs":300,
        "seed":42})

%time glm_grid.train(x = train_X, y = train_y, training_frame = train,
validation_frame = valid)

glm Grid Build progress: |  

██████████████████████████████████████████████████████| (done) 100%
CPU times: total: 1.61 s  

Wall time: 5min 2s
```

Hyper-Parameter Search Summary: ordered by increasing logloss

alpha	missing_values_handling	model_ids	logloss
0.0	Skip		
		glm random grid model 11	0.39160808622118365

0.0	MeanImputation
glm_random_grid_model_34	0.39160808622118365
0.44	MeanImputation
glm_random_grid_model_28	0.3917569991015515
0.48	Skip
glm_random_grid_model_2	0.3917586027842902
0.48	MeanImputation
glm_random_grid_model_36	0.3917586027842902
0.43	MeanImputation
glm_random_grid_model_10	0.39176212599720756
0.43	Skip
glm_random_grid_model_22	0.39176212599720756
0.47000000000000003	MeanImputation
glm_random_grid_model_26	0.3917631576492127
0.37	Skip
glm_random_grid_model_13	0.3917744343513941
0.25	Skip
glm_random_grid_model_19	0.3917771796917925
---	---
---	---

0.07	MeanImputation
glm_random_grid_model_20	0.39183529584589244
0.07	Skip
glm_random_grid_model_21	0.39183529584589244
0.11	MeanImputation
glm_random_grid_model_25	0.3918400598420588
0.1	MeanImputation
glm_random_grid_model_24	0.39184069991759385
0.08	Skip
glm_random_grid_model_23	0.39185013254495743
0.06	MeanImputation
glm_random_grid_model_9	0.3918788835170476
0.05	MeanImputation
glm_random_grid_model_4	0.3919370150415301
0.03	Skip
glm_random_grid_model_30	0.39213545399292443
0.02	Skip
glm_random_grid_model_27	0.3923318174271338
0.29	Skip
glm_random_grid_model_37	0.39418362265362783

[37 rows x 5 columns]

```
sorted_glm_grid = glm_grid.get_grid(sort_by = 'aucpr', decreasing =
True)
sorted_glm_grid.sorted_metric_table()
best_model_id = sorted_glm_grid.sorted_metric_table()['model_ids'][0]
best_model_id
'glm_random_grid_model_11'
```

he grid search results sorted by the AUC-PR metric in decreasing order. AUC-PR (Area Under the Precision-Recall Curve) is a metric commonly used to evaluate the performance of binary classification models, especially when dealing with imbalanced datasets. Sorting in decreasing order means that the models with the highest AUC-PR values will appear at the top of the sorted list. Higher AUC-PR values indicate better precision-recall trade-offs in the models.

sorted_metric_table() function provides an easy way to examine and analyze the results, allowing you to identify the best-performing models based on the chosen metric (AUC-PR in this case).

As in this case we can see the best model is glm_random_grid_model_11 with AUCPR values as 0.439303

```
tuned_glm = sorted_glm_grid.models[0]
tuned_glm.summary()

GLM Model: summary
  family    link    regularization    lambda_search
number_of_predictors_total    number_of_active_predictors
number_of_iterations    training_frame
--  -----  -----  -----
-----
      binomial  logit  Ridge ( lambda = 1.252E-5 )  nlambda = 30,
lambda.max = 12.525, lambda.min = 1.252E-5, lambda.1se = -1.0  27
27                                54                                py_17_sid_b6e5
```

Here we receive the best model for grid search along with the parameters from the best model

Lets evaluate the model on the validation set

```
tuned_glm_perf = tuned_glm.model_performance(valid)
print("Default GLM AUCPR: %.4f \nTuned GLM AUCPR: %.4f" % (glm.aucpr(),
tuned_glm_perf.aucpr()))

Default GLM AUCPR: 0.4455
Tuned GLM AUCPR: 0.4393
```

Which suggests are default GLM model having higher AUCPR generalizes better than the tuned model

As such we will do Scoring in Scikit-Learn Model which has a better threshold

Saving the H2o model in the Artifacts

```
# Define the directory path and the model file name separately
model_directory =
"D:/Work/Gre/UTD/Courses/Fall/MIS6341/Softwares/Python/ml-fall-2023/
Project1/artifacts_h2o"
model_filename = "glm_model_h2o.pkl"

# Get the H2O model by its ID
model_h2o = h2o.get_model(best_model_id)

# Construct the full model path
model_path = f"{model_directory}/{model_filename}"

# Create an artifacts dictionary and include the model file path
artifacts_dict = {}
artifacts_dict["h2o_model_path"] = model_path

# Save the H2O model to the specified file path
h2o.save_model(model_h2o, model_path)

'D:\\Work\\Gre\\UTD\\Courses\\Fall\\MIS6341\\Softwares\\Python\\ml-
fall-2023\\Project1\\artifacts_h2o\\glm_model_h2o.pkl\\
glm_random_grid_model_11'

def score_with_h2o_model(input_data):
    import h2o
    try:
        h2o.cluster().shutdown()
    except:
        pass
    from h2o.frame import H2OFrame
    h2o.init(max_mem_size = "4G", nthreads=16)
    try:
        # Load the saved H2O model
        model_path =
"D:/Work/Gre/UTD/Courses/Fall\\MIS6341/Softwares/Python/ml-fall-2023/
Project1/artifacts_h2o/glm_model_h2o.pkl/glm_random_grid_model_11"
        loaded_model = h2o.load_model(model_path)

        # Convert input_data to an H2O frame
        input_h2o = h2o.H2OFrame(input_data)

        # Use the loaded model for scoring
        predictions = loaded_model.predict(input_h2o)
```



```

# Extract the predictions as a Pandas DataFrame
predictions_df = predictions.as_data_frame()

# Return the prediction results
return predictions_df

except Exception as e:
    return f"Error: {e}"

score_with_h2o_model(valid)

H2O session _sid_b6e5 closed.
Checking whether there is an H2O instance running at
http://localhost:54321.

.... not found.
Attempting to start a local H2O server...
; Java HotSpot(TM) 64-Bit Server VM (build 25.361-b09, mixed mode)
Starting server from D:\Work\Gre\UTD\Courses\Fall\MIS6341\Softwares\
Python\ml-fall-2023\Lib\site-packages\h2o\backend\bin\h2o.jar
Ice root: C:\Users\Asus\AppData\Local\Temp\tmps874sy9h
JVM stdout: C:\Users\Asus\AppData\Local\Temp\tmps874sy9h\
h2o_Asus_started_from_python.out
JVM stderr: C:\Users\Asus\AppData\Local\Temp\tmps874sy9h\
h2o_Asus_started_from_python.err
Server is running at http://127.0.0.1:54321
Connecting to H2O server at http://127.0.0.1:54321 ... successful.

-----
H2O_cluster_uptime:      04 secs
H2O_cluster_timezone:    America/Chicago
H2O_data_parsing_timezone: UTC
H2O_cluster_version:     3.42.0.3
H2O_cluster_version_age: 2 months and 14 days
H2O_cluster_name:        H2O_from_python_Asus_vzln0q
H2O_cluster_total_nodes: 1
H2O_cluster_free_memory: 3.548 Gb
H2O_cluster_total_cores: 16
H2O_cluster_allowed_cores: 16
H2O_cluster_status:      locked, healthy
H2O_connection_url:       http://127.0.0.1:54321
H2O_connection_proxy:     {"http": null, "https": null}
H2O_internal_security:    False
Python_version:           3.10.11 final
-----

'Error: Argument `python_obj` should be a None | list | tuple | dict |
numpy.ndarray | pandas.DataFrame | scipy.sparse.issparse, got H2OFrame
City      State      Zip      Bank      BankState      NAICS      NoEmp
NewExist  CreateJob    RetainedJob    FranchiseCode    UrbanRural

```

RevLineCr	LowDoc	DisbursementGross	BalanceGross	GrAppv
SBA_Appv	MIS_Status	Log_DisbursementGross	Log_GrAppv	
Log_SBA_Appv	Log_BalanceGross	TotalJobs	IncomeToLoanRatio	
EmployeesToLoanRatio	JobPerLoan	Gauren_SBA_Appv	DefaultRate\	
n0.11465	0.184773	93001	0.0314465	0.218517 235910
0.600407	0.17044	-0.0353733	-0.0454543	1
0.0716743	0.15307	0.187063	0.358949	-0.00229552
0.394801	0.410973	0	0.306712	
0.332752	0.344279	-0.00229816	-0.0808276	
0.873414		1.46094	-0.196674	0.960651
17.5096\n0.337588	0.197662	11225	0.272221	0.220168 621111
-0.140136	0.186978	-0.022536	-0.0326467	0
0.243491	0.15307	0.187063	-0.451437	-0.00229552
-0.429073	-0.497867	0	-0.600454	-
0.560494	-0.68889	-0.00229816	-0.0551827	
0.906743		0.281473	0.110838	0.861822
17.5096\n0.0615385	0.124634	54935	0.175694	0.117429 453220
-0.140136	0.186978	-0.0353733	-0.0454543	1
0.0716743	0.15307	0.0897581	-0.593428	-0.00229552
-0.573427	-0.547995	0	-0.899994	-
0.851971	-0.794062	-0.00229816	-0.0808276	
1.08291		0.255725	0.147497	1.04641
17.5096\n0.185864	0.188249	77379	0.147425	0.139976 332996
0.788908	0.17044	-0.0353733	-0.0454543	1
0.0716743	0.15307	0.187063	2.14249	-0.00229552
2.20803	2.02814	0	1.14502	
1.16566	1.10795	-0.00229816	-0.0808276	
1.05638		0.388981	-0.039853	1.0887
17.5096\n0.275041	0.184773	90001	0	0.218517 448120
0.654264	0.17044	0.0288134	-0.0454543	1
0.243491	0.15307	0.187063	3.9295	-0.00229552
4.24307	5.44119	0	1.59524	
1.65691	1.86271	-0.00229816	-0.0166409	
0.722176		0.120243	-0.00305832	0.779804
17.5096\n0.189995	0.156142	64068	0.145683	0.126444 811192
-0.140136	0.17044	-0.0353733	-0.0454543	1
0.0716743	0.15307	0.187063	-0.143213	-0.00229552
-0.115719	-0.129536	0	-0.154566	-
0.12298	-0.138729	-0.00229816	-0.0808276	
1.10558		1.08183	0.623976	0.893333
17.5096\n0.175011	0.124634	54455	0.116923	0.117429 518210
-0.140136	0.17044	-0.0353733	-0.0411851	0
0.187265	0.251569	0.187063	-0.464802	-0.00229552
-0.573427	-0.587225	0	-0.625118	-
0.851971	-0.884854	-0.00229816	-0.0765584	
0.791522		0.238641	0.130373	0.976502
17.5096\n0.117995	0.140225	98116	0.175694	0.159167 812113
-0.072814	0.17044	-0.0310942	-0.01557	0
0.243491	0.251569	0.187063	-0.610744	-0.00229552

```
-0.591031  -0.598123      0      -0.943518  -
0.894116   -0.911609     -0.00229816  -0.0466641
1.0211      0.121737     0.0780177      0.988143
17.5096\n0.115      0.12935      2745      0.0896552      0.140419      448130
-0.140136   0.17044      -0.0353733      -0.0411851      1
0.243491    0.251569     0.187063      -0.489532      -0.00229552
-0.467802   -0.521841      0      -0.672428      -
0.63074     -0.737812     -0.00229816  -0.0765584
0.938087      0.268541     0.146708      0.896445
17.5096\n0.276347   0.224201     60657     0.175694      0.159167      812111
0.0214369   0.17044      -0.0353733      -0.0454543      1
0.243491    0.15307     0.187063      -0.181308      -0.00229552
-0.154448   -0.100549      0      -0.200048      -
0.167766    -0.105971     -0.00229816  -0.0808276
1.80318      -0.213198     0.803859      1.53604
17.5096\n[84951 rows x 29 columns]\n'
```